

# Terraform: Infrastructure as Code

---

Martin Schütte

20 August 2017




# Concepts

# From Servers ...



by Rodzilla at Wikimedia Commons (CC-BY-SA-3.0)



## Services have APIs

- Starting servers is just a command line or function call
  - Add to build process (phoenix/immutable servers)
  - Replace “click paths” with source code in VCS
  - Fewer “black box” setup steps, better team handovers
- ⇒ Infrastructure as Code

# Services also need Configuration Management

- Lifecycle awareness, not just a `setup.sh`
  - Multiple stages/environments
  - Specification, documentation, policy enforcement
- ⇒ Tool support



# TERRAFORM

Build, Combine, and Launch Infrastructure



## Example: Simple Webservice (part 1)

```
### AWS Setup
provider "aws" {
    profile = "${var.aws_profile}"
    region  = "${var.aws_region}"
}

# Queue
resource "aws_sqs_queue" "importqueue" {
    name = "${var.app_name}-${var.aws_region}-importqueue"
}

# Storage
resource "aws_s3_bucket" "importdisk" {
    bucket = "${var.app_name}-${var.aws_region}-importdisk"
    acl    = "private"
}
```

## Example: Simple Webservice (part 2)

```
### Heroku Setup
provider "heroku" { ... }

# Importer
resource "heroku_app" "importer" {
  name      = "${var.app_name}-${var.aws_region}-import"
  region    = "eu"
  config_vars {
    SQS_QUEUE_URL = "${aws_sqs_queue.importqueue.id}"
    S3_BUCKET      = "${aws_s3_bucket.importdisk.id}"
  }
}

resource "heroku_addon" "mongolab" {
  app    = "${heroku_app.importer.name}"
  plan   = "mongolab:sandbox"
}
```

# Core Ideas in Terraform

- Simple model of resource entities with attributes
- Stateful lifecycle with CRUD operations
- Declarative configuration
- Dependencies by inference
- Parallel execution

# Core Concepts in Terraform

- Provider: a source of resources  
(usually with an API endpoint & authentication)
- Resource: every thing “that has a set of configurable attributes and a lifecycle (create, read, update, delete)” – implies ID and state
- Data Source: information read from provider  
(e.g. lookup own account ID or AMI-ID)
- Provisioner: initialize a resource with local or remote scripts

# Design Choices in Terraform

- Order: directed acyclic graph of all resources
- Plan: generate an execution plan for review before applying a configuration
- State: execution result is kept in state file (local or remote)
- Lightweight: little provider knowledge, no error handling

# Available services

Providers:	Resources:	Provisioners:
<ul style="list-style-type: none"><li>• AWS</li><li>• Azure</li><li>• Google Cloud</li><li>• Alicloud</li><li>• Heroku</li><li>• DNSMadeEasy</li><li>• OpenStack</li><li>• Docker</li><li>• ...</li></ul>	<ul style="list-style-type: none"><li>• aws_instance</li><li>• aws_vpc</li><li>• aws_iam_user</li><li>• azurerm_subnet</li><li>• azurerm_dns_zone</li><li>• azure_instance</li><li>• aws_iam_user</li><li>• heroku_app</li><li>• postgresql_schema</li><li>• ...</li></ul>	<ul style="list-style-type: none"><li>• chef</li><li>• file</li><li>• local-exec</li><li>• remote-exec</li></ul>

- Hashicorp Configuration Language (HCL),  
think “JSON-like but human-friendly”
- Variables
- Interpolation, e.g.  
`"number ${count.index + 1}"`
- Attribute access with `resource_type.resource_name`
- Few build-in functions, e.g.  
`base64encode(string), format(format, args...)`

# HCL vs. JSON


```
# An AMI
variable "ami" {
    description = "custom AMI"
}

/* A multi
   line comment. */
resource "aws_instance" "web" {
    ami = "${var.ami}"
    count = 2
    source_dest_check = false


    connection {
        user = "root"
    }
}
```

```
{
    "variable": {
        "ami": {
            "description": "custom AMI"
        }
    },
    "resource": {
        "aws_instance": {
            "web": {
                "ami": "${var.ami}",
                "count": 2,
                "source_dest_check": false,
                "connection": {
                    "user": "root"
                }
            }
        }
    }
}
```

# terraform graph | dot -Tpdf



# Terraform Process



# Example: Add Provisioning

```
# Importer
resource "heroku_app" "importer" {
  name    = "${var.app_name}-${var.aws_region}-import"
  region  = "eu"

  config_vars { ... }

  provisioner "local-exec" {
    command = <<EOT
cd ~/projects/go-testserver &&
git remote add heroku ${heroku_app.importer.git_url} &&
git push heroku master
EOT
  }
}
```

## Example: Add Outputs

```
# Storage
resource "aws_s3_bucket" "importdisk" { ... }

# Importer
resource "heroku_app" "importer" { ... }

# Outputs
output "importer_bucket_arn" {
    value = "${aws_s3_bucket.importdisk.arn}"
}

output "importer_url" {
    value = "${heroku_app.importer.web_url}"
}

output "importer_gitrepo" {
    value = "${heroku_app.importer.git_url}"
}
```

## Example: Add Lifecycle Meta-Parameter

```
# Storage
resource "aws_s3_bucket" "importdisk" {
  bucket = "${var.app_name}-${var.aws_region}-importdisk"
  acl    = "private"

  lifecycle {
    prevent_destroy = true
  }
}
```

# Demo

```
$ terraform init  
$ terraform validate  
$ terraform plan -out=my.plan  
$ terraform show my.plan  
$ terraform apply my.plan  
  
$ terraform output  
$ terraform output -json  
$ terraform output importer_url  
$ curl -s $(terraform output importer_url)  
  
$ terraform graph | dot -Tpdf > graph.pdf && evince graph.pdf  
  
$ terraform plan -destroy  
$ terraform destroy
```

# Features

“Plain terraform code” lacks structure and reusability

## Modules

- are subdirectories with self-contained terraform code
- may be sourced from Git, Mercurial, HTTPS locations
- use variables and outputs to pass data

# Example Module

```
module "database" {
  source = "github.com/terraform-community-modules/tf_aws_rds"

  # DB Instance Inputs
  rds_instance_identifier = "${terraform.workspace}-${var.app}-db"
  rds_allocated_storage   = "${var.database_size}"

  database_name      = "${var.database_name}"
  database_user       = "${var.database_user}"
  database_password  = "${var.database_password}"

  # DB Subnet Inputs
  subnets = ["${aws_subnet.dbnet.*.id}"]
  rds_vpc_id = "${data.aws_vpc.app.id}"

  tags {
    Name = "${terraform.workspace} - ${var.app} - DB"
  }
}
```

- Terraform keeps known state of resources
- Defaults to local state in `terraform.tfstate`
- Optional remote state with different backends  
(S3, Azure Storage, Consul, ...)
  - Useful to sync multiple team members
  - May need additional mutex mechanism  
(v0.9 added state locking for Local, S3, and Consul)
  - Remote state is a data source

# Example: Using State Import

```
$ terraform import azurerm_storage_account.my_storage_account \
/subscriptions/e9b2ec19-ab6e-4547-a3ec-5a58e234ce5e/resourceGroups/
demo-res-group/providers/Microsoft.Storage/storageAccounts/demostorage20170418

azurerm_storage_account.my_storage_account: Importing from ID ...
azurerm_storage_account.my_storage_account: Import complete!
  Imported azurerm_storage_account (ID: ...)
azurerm_storage_account.my_storage_account: Refreshing state... (ID: ...)
```

Import success! The resources imported are shown above. These are now in your Terraform state. Import does not currently generate configuration, so you must do this next. If you do not create configuration for the above resources, then the next 'terraform plan' will mark them for destruction.

```
$ terraform state list
azurerm_storage_account.my_storage_account

$ terraform state show azurerm_storage_account.my_storage_account
id                  = /subscriptions/e9b2ec19...
account_kind        = Storage
account_type        = Standard_LRS
location            = westeurope
name                = demostorage20170418
...
...
```

## Example: Use Remote State (with Workspaces)

```
terraform {
  required_version = ">= 0.9.8"
  environment = "${terraform.workspace}"
  backend "s3" {
    bucket  = "ms-terraform-state"
    key     = "infra/ms-tf-demo/state"
    region  = "eu-central-1"
  }
}

$ terraform workspace new prod
$ terraform workspace new dev
$ aws s3 ls --recursive "s3://ms-terraform-state/"
... 282 workspace:/dev/infra/ms-tf-demo/state
... 282 workspace:/prod/infra/ms-tf-demo/state
```

## Example: Use Remote State to Chain Projects

```
data "terraform_remote_state" "infra" {
  backend = "s3"
  config {
    bucket    = "ms-terraform-state"
    key       = "workspace:/${terraform.workspace}/infra/_
    ↪ ms-tf-demo/state"
    region   = "eu-central-1"
  }
}

resource "aws_instance" "foo" {
  # use state from vpc_project
  subnet_id      =
  ↪ "${data.terraform_remote_state.infra.app_subnet_id}"
  instance_type  = "t2.micro"
  ami            = "ami-b968bad6"
}
```

## Example: Using Data Source to Lookup Data

```
# searches for most recent tagged AMI in own account
data "aws_ami" "webami" {
    most_recent = true
    owners      = ["self"]

    filter {
        name    = "tag:my_key"
        values = ["my_value"]
    }
}

# use AMI
resource "aws_instance" "web" {
    instance_type = "t2.micro"
    ami           = "${data.aws_ami.webami.id}"
}
```

## Example: “External” Data Source

```
data "external" "dyndns" {
  program = ["bash", "${path.module}/variomedia_dyndns.sh"]

  query = {
    hostname  = "aws-demo.martin-schuette.de"
    ipaddress = "${aws_eip.foo.public_ip}"
  }
}
```

# How to Write Own Plugins

Now:

- Learn you some Golang 
- Use the schema helper lib
- Adapt to model of  
Provider (setup steps, authentication) and  
Resources (arguments/attributes and CRUD methods)
- Start reading of simple plugins like  
`builtin/providers/mysql` 

Future:

-  interface, support for Python, Ruby, C#, Java, ...

# Usage

## General Problemes for all Tools

---

- Testing is inherently difficult
- Provider coverage largely depends on community
- Resource model mismatches, e.g. with Heroku apps
- Ignorant of API rate limits, account ressource limits, etc.

Under active development,  
current version 0.10.2 (August 16)

- Modules are very simple
- Lacking syntactic sugar  
(e.g. aggregations, common repetitions)
- Big improvements in state management
- Large variation in provider support, new project boundaries

## Recent Features in 0.7–0.10

- State Import
- Data Sources
- Workspaces (previously: State Environments)
- Separate sub-projects for providers  
[terraform-providers](#) 
- Support for gRPC-based plugins, i. e. providers in other languages

# Comparable Tools

## Configuration Management Tools:

- SaltStack Salt Cloud
- Ansible modules
- Puppet modules

## Vendor Tools:

- Azure Resource Manager Templates
- AWS CloudFormation
- OpenStack Heat

- Avoid user credentials in Terraform code,  
use e.g. profiles and *assume-role* wrapper scripts
- At least use separate user credentials,  
know how to revoke them
- To hold credentials in VCS use PGP encryption,  
e.g. with [Blackbox](#)

## Workflow (contd.)


- Use a VCS, i. e. git
- Namespaces! – Always add some  
"\${var.shortname}-\${var.env}"
  - per project
  - per region
  - per account
  - per provider
- Use remote state and consider access locking,  
e.g. with a single build server
- Take a look at [Hashicorp Atlas](#) and its workflow

# Hashicorp Toolset

## PROVISION, SECURE, AND RUN

ANY INFRASTRUCTURE FOR ANY APPLICATION

 LEARN THE HASHICORP SUITE >



Seven elements of the modern Application Lifecycle



# Links and Resources

- [Terraform.io](#) and [hashicorp/terraform](#) 
- [terraform-providers](#) 
- [terraform-community-modules](#) 
- [newcontext/kitchen-terraform](#) 
- Terraforming – Export existing AWS resources
- Terraform: Beyond the Basics with AWS
- A Comprehensive Guide to Terraform
- Terraform, VPC, and why you want a tfstate file per env

# Books

*Hopefully, deployments will become routine and boring—and in the world of operations, boring is a very good thing.*

— *Terraform: Up & Running* by Yevgeniy Brikman



*Defining system infrastructure as code and building it with tools doesn't make the quality any better. At worst, it can complicate things.*

— *Infrastructure as Code* by Kief Morris

Thank You! – Questions?

**Workshop**  
**Terraform und AWS**  
at 14:00 h in C 120



Martin Schütte  
@m\_schuett  
[info@martin-schuette.de](mailto:info@martin-schuette.de)

[slideshare.net/mschuett/](https://slideshare.net/mschuett/)   
[tinyurl.com/froscon17-tf](https://tinyurl.com/froscon17-tf)